




Harbin Institute of Technology, Shenzhen



Last Time


- What is machine learning?
- Linear Classifier
 - Logistic Regression
 - Support Vector Machine
- Structured Learning
 - Hidden Markov Model
 - Condition Random Field





Today's class

- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





From ML to DL

- Let's recap machine learning.
- **A machine learning algorithm:**
 - **input:** Training Data
 - “learns” from the training data.
 - **output:** A “prediction function” that produces output y given input x .



From ML to DL

- Let's recap machine learning.
- e.g. Logistic Regression


$$p(y = 1|\mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x} + b) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x} + b)}$$
$$p(y = 0|\mathbf{x}) = 1 - \sigma(\mathbf{w}^\top \mathbf{x} + b)$$



From ML to DL

- e.g. Logistic Regression


Var	Definition
x_1	count(positive lexicon) $\in$ doc
x_2	count(negative lexicon) $\in$ doc
x_3	$\begin{cases} 1 & \text{if "no" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$
x_4	count(1st and 2nd pronouns $\in$ doc)
x_5	$\begin{cases} 1 & \text{if "!" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$
x_6	log(word count of doc)



$$\mathbf{x} = [3, 2, 1, 3, 0, 4.19]^T$$




From ML to DL



- **Preprocessing:** remove noisy text, remove stop words
- **Feature Creation:** create features from original data
- **Feature Transform:** feature processing, i.e. dimensionality reduction
 - feature extraction: create a subset of new features by combinations of the existing features. i.e. PCA, LDA
 - feature selection: choose a subset of all the features (the ones more informative). i.e. mutual information, TF-IDF





From ML to DL



- The performance of ML model tends to be similar.
- **Feature Engineering is more important!**




From ML to DL





From ML to DL



$$f(x) = \sigma(\mathbf{w}^\top \mathbf{x})$$


stronger expression ability
layer by layer feature learning




From ML to DL

- Three components
 - Representation: neural network.
 - Evaluation: cross-entropy loss, etc.
 - Optimization: stochastic gradient descent.

What others think DL is...




What it actually is ...





From ML to DL


- Why DL didn't work before?
 - large amounts of data
 - cheaper and faster computers
 - advanced ML techniques





Today's class


- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





Biological Neuron

- 生物神经元





A neuron has only two states:
Excitation and Inhibition



Artificial Neuron

- 人工神经元



$$\sigma(\mathbf{w}^\top \mathbf{x} + b) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x} + b)}$$




Artificial Neural Network

- An ANN is based on a collection of directed connected neurons.
- Three problems:
 - **neuron activations rules:** mapping from input to output, usually non-linear function
 - **network topology:** Connection between different neurons
 - **learning algorithm:** Learning appropriate parameters by training




Artificial Neural Network


- Parallel distributed structure
 - three basic types of neural network



(a) Feed Forward Network



(b) Memory Network




(c) Graph Network

- but actually, most networks are **composite network** that consists of different type of networks.



Feed Forward Neural Network

- A feedforward neural network is an artificial neural network wherein connections between the nodes do not form a cycle.





Universal Approximation Theorem

- A feed forward network with a single hidden layer containing a finite number of neurons can approximate continuous functions on compact subsets of $\mathbf{R}^n$, under mild assumptions on the activation function.
- The theorem thus states that simple neural networks can represent a wide variety of interesting functions when given appropriate parameters.

A two-layer neural network can simulate any function



Popular Deep Learning Framework

- Easy and fast prototyping
- Automatic gradient computation
- Support CPUs and GPUs



 PaddlePaddle






Three Steps of Deep Learning

- Deep Learning is so simple...

- ① design a network
- ② design a loss function
- ③ optimize the parameters




Deep Learning is easy :-)





Today's class


- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





Language Modeling

- Task: Language Modeling
 - to predict what word comes next.




- Formally, given a sequence of words $x^{(1)}, x^{(2)}, \dots, x^{(t)}$, compute the probability distribution of the next words $x^{(t+1)}$.

$$P(x^{(t+1)} | x^{(t)}, \dots, x^{(1)})$$



Language Modeling






Language Modeling

- A fixed-window neural Language Model

- fixed window, e.g. $n_{window} = 4$
- parameters: W , U


- Weakness

- fixed window is **too small**
- enlarging window enlarges W
- the interaction between $x^{(1)}$ and $x^{(2)}$ is not modeled.





Recurrent Neural Networks




Core idea: Apply the same weights W repeatedly.



Recurrent Neural Networks

hidden states

$$h^{(t)} = \sigma(W_h h^{(t-1)} + W_e e^{(t)} + b_1)$$






Recurrent Neural Networks

- Advantages:


- can process **any length** input
- computation for step t can (in theory) use information from **many steps back**
- **model size doesn't increase** for longer input
- same weights applied on every timestep, so there is **symmetry** in how inputs are processed





RNN can be used for tagging


- e.g. Part-of-speech tagging





RNN can be used for sentence classification


- e.g. sentiment classification





RNN can be used for sentence classification


- e.g. sentiment classification





RNN can be used for sentence classification


- e.g. sentiment classification





RNN can be used to generate text

- e.g. speech recognition, machine translation, summarization






Long-Term Dependencies Problem

- Disadvantages of RNN:
 - memory capacity: as $h^{(t)}$ continues to accumulate new input information, saturation will occur.

as the *proctor* started the clock, the students opened their *exam*.

当 监考员 宣布 开考, 学生 开始 了 考试”





Long Short Term Memory (LSTM)

- On step t , there is a **hidden state** $h^{(t)}$ and a **cell state** $c^{(t)}$
 - Both are vectors length n
 - The cell stores **long-term information**
 - The LSTM can **erase**, **write** and **read** information from the cell
- The selection of which information is erased/written/read is controlled by three corresponding **gates**
 - The gates are also vectors length n
 - On each timestep, each element of the gates can be **open** (1), **closed** (0), or somewhere in-between.
 - The gates are **dynamic**: their value is computed based on the current context



Long Short Term Memory (LSTM)

We have a sequence of inputs $\mathbf{x}^{(t)}$, and we will compute a sequence of hidden states $\mathbf{h}^{(t)}$ and cell states $\mathbf{c}^{(t)}$. On timestep t :

Forget gate: controls what is kept vs forgotten, from previous cell state

Input gate: controls what parts of the new cell content are written to cell

Output gate: controls what parts of cell are output to hidden state

New cell content: this is the new content to be written to the cell

Cell state: erase (“forget”) some content from last cell state, and write (“input”) some new cell content

Hidden state: read (“output”) some content from the cell

Sigmoid function: all gate values are between 0 and 1

$$\mathbf{f}^{(t)} = \sigma \left(\mathbf{W}_f \mathbf{h}^{(t-1)} + \mathbf{U}_f \mathbf{x}^{(t)} + \mathbf{b}_f \right)$$

$$\mathbf{i}^{(t)} = \sigma \left(\mathbf{W}_i \mathbf{h}^{(t-1)} + \mathbf{U}_i \mathbf{x}^{(t)} + \mathbf{b}_i \right)$$

$$\mathbf{o}^{(t)} = \sigma \left(\mathbf{W}_o \mathbf{h}^{(t-1)} + \mathbf{U}_o \mathbf{x}^{(t)} + \mathbf{b}_o \right)$$

$$\tilde{\mathbf{c}}^{(t)} = \tanh \left(\mathbf{W}_c \mathbf{h}^{(t-1)} + \mathbf{U}_c \mathbf{x}^{(t)} + \mathbf{b}_c \right)$$

$$\mathbf{c}^{(t)} = \mathbf{f}^{(t)} \circ \mathbf{c}^{(t-1)} + \mathbf{i}^{(t)} \circ \tilde{\mathbf{c}}^{(t)}$$

$$\mathbf{h}^{(t)} = \mathbf{o}^{(t)} \circ \tanh \mathbf{c}^{(t)}$$



Gates are applied using element-wise product

All these are vectors of same length n



Long Short Term Memory (LSTM)


You can think of the LSTM equations visually like this:





Long Short Term Memory (LSTM)

You can think of the LSTM equations visually like this:





Long Short Term Memory (LSTM)

- The LSTM architecture makes it easier for the RNN to **preserve information over many timesteps**.
- In 2013-2015, LSTMs started achieving state-of-the-art results
 - successful tasks include: **handwriting recognition, speech recognition, machine translation, parsing, image captioning**
 - LSTM became the dominant approach



Gated Recurrent Units (GRU)

- Proposed by Cho et al. in 2014 as a simpler alternative to the LSTM.
- On each timestep t we have input $\mathbf{x}^{(t)}$ and hidden state $\mathbf{h}^{(t)}$ (no cell state).

Update gate: controls what parts of hidden state are updated vs preserved

$$\mathbf{u}^{(t)} = \sigma \left(\mathbf{W}_u \mathbf{h}^{(t-1)} + \mathbf{U}_u \mathbf{x}^{(t)} + \mathbf{b}_u \right)$$

Reset gate: controls what parts of previous hidden state are used to compute new content

$$\mathbf{r}^{(t)} = \sigma \left(\mathbf{W}_r \mathbf{h}^{(t-1)} + \mathbf{U}_r \mathbf{x}^{(t)} + \mathbf{b}_r \right)$$

New hidden state content: reset gate selects useful parts of prev hidden state. Use this and current input to compute new hidden content.

$$\tilde{\mathbf{h}}^{(t)} = \tanh \left(\mathbf{W}_h (\mathbf{r}^{(t)} \circ \mathbf{h}^{(t-1)}) + \mathbf{U}_h \mathbf{x}^{(t)} + \mathbf{b}_h \right)$$

$$\mathbf{h}^{(t)} = (1 - \mathbf{u}^{(t)}) \circ \mathbf{h}^{(t-1)} + \mathbf{u}^{(t)} \circ \tilde{\mathbf{h}}^{(t)}$$

Hidden state: update gate simultaneously controls what is kept from previous hidden state, and what is updated to new hidden state content

How does this solve vanishing gradient?

Like LSTM, GRU makes it easier to retain info long-term (e.g. by setting update gate to 0)




LSTM vs GRU

- Researchers have proposed many gated RNN variants, but LSTM and GRU are the most widely-used
- Rule of thumb: LSTM is a good default choice (especially if your data has particularly long dependencies, or you have lots of training data); Switch to GRUs for speed and fewer parameters.



Bidirectional RNNs





Bidirectional RNNs

On timestep t :

This is a general notation to mean “compute one forward step of the RNN” – it could be a vanilla, LSTM or GRU computation.

Forward RNN $\vec{h}^{(t)} = \text{RNN}_{\text{FW}}(\vec{h}^{(t-1)}, \mathbf{x}^{(t)})$

Backward RNN $\overleftarrow{h}^{(t)} = \text{RNN}_{\text{BW}}(\overleftarrow{h}^{(t+1)}, \mathbf{x}^{(t)})$

Generally, these two RNNs have separate weights

Concatenated hidden states $\mathbf{h}^{(t)} = [\vec{h}^{(t)}; \overleftarrow{h}^{(t)}]$


We regard this as “the hidden state” of a bidirectional RNN. This is what we pass on to the next parts of the network.



Bidirectional RNNs

Multi-layer RNNs


The hidden states from RNN layer i are the inputs to RNN layer $i+1$





Today's class

- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





Convolutional Neural Models

- Image Classification

Image Classification



64x64

→ Cat? (0/1)



Convolutional Neural Models

- Image Classification by FNN

- input size=64*64*3=12288


- let hidden size=1024

- $W \in R^{12288 \times 1024}$

- two much parameters!



64x64





Convolutional Neural Models

- Motivation





Convolutional Neural Models


- Motivation






Convolutional Neural Models

- Motivation



64x64





Convolutional Neural Models

- Convolution

1	1	1	1	1
		$\times -1$	$\times 0$	$\times 0$
-1	0	-3	0	1
		$\times 0$	$\times 0$	$\times 0$
2	1	1	-1	0
		$\times 0$	$\times 0$	$\times 1$
0	-1	1	2	1
1	2	1	1	1

 $\otimes$

1	0	0
0	0	0
0	0	-1


 $=$

0	-2	-1
2	2	4
-1	0	0




Convolutional Neural Models


- Convolution




stride(步长)=1, padding(填充)=0



stride=2, padding=0



stride=1, padding=1




stride=2, padding=1




Convolutional Neural Models

- Replace fully-connected layer with convolution layer



(a) 全连接层




(b) 卷积层



Convolutional Neural Models

- Multiple filter





Convolutional Neural Models

- Pooling Layer

1	1	2	0
0	1	1	2
0	0	1	3
0	0	1	1

输入特征映射 X^d

max pooling
→


1	2
0	3

输出特征映射 Y^d



Convolutional Neural Models


- CNN for NLP





Convolutional Neural Models

- TextCNN





Convolutional Neural Models

- TextCNN




Figure 1: Model architecture with two channels for an example sentence.



Convolutional Neural Models

- Why does textCNN work?

not recommend”), but not at the document level. Indeed, in trying to determine the polarity of a movie review, we don’t really care whether “not bad, quite good” is found at the start or at the end of the document. We just need to capture the fact that “not” precedes “bad”, and so forth. Note that CNNs are not able to encode long-range dependencies, and therefore, for some tasks like language modeling, where long-distance dependence matters, recurrent architectures such as LSTMs are preferred.



Today's class

- From ML to DL
- Neuron and ANN

- DL models


- RNN 循环神经网络

- CNN 卷积神经网络

- Encoder-Decoder 编码器-解码器网络

- Attention 注意力机制

- Transformer 网络





Encoder-Decoder

- Sequence to sequence

- core component


- RNN (GRU、LSTM)


- extension

- with attention mechanism
 - with copy mechanism
 -


- decoding strategy

- Greedy decoding
 - Beam search
 -





-






Encoder-Decoder

- Drawbacks


One of the main drawbacks of this network is its **inability to extract strong contextual relations from long semantic sentences**, that is if a particular piece of long text has some context or relations within its substrings, then a basic seq2seq model cannot identify those contexts and therefore, somewhat decreases the performance of the model and eventually, decreasing accuracy.





Today's class

- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





Sentiment Classification

- Task: Aspect-based Sentiment classification
 - given a review and a aspect term, to determine the sentiment polarity expressed by the review on this aspect term.


"The restaurant was expensive, but the menu was great" → {price: negative, food: positive}



Sentiment Classification

- Baseline: LSTM
 - construct a new sentence
 - “Aspect food. The restaurant was expensive, but the menu was great”
 - input this sentence to LSTM
 - perform max-pooling

$$\mathbf{h} = \text{maxpooling}(\mathbf{h}^{(1)}, \dots, \mathbf{h}^{(T)})$$



“The restaurant was expensive, but the menu was great” → {price: negative, food: positive}



Sentiment Classification

- Attention-based LSTM
 - input original sentence to LSTM
 - *“The restaurant was expensive, but the menu was great”*
 - weighted sum

$$\text{h} = \sum_{i=1}^T \alpha_i \mathbf{h}^{(i)}$$

- large weight for *“menu”* and *“great”*
- Small weight for *“restaurant”*, *“expensive”* and *“the”*



Sentiment Classification

- Attention-based LSTM
 - weighted sum

$$\cancel{\mathbf{h} = \text{maxpooling}(\mathbf{h}^{(1)}, \dots, \mathbf{h}^{(T)})} \quad \mathbf{h} = \sum_{i=1}^T \alpha_i \mathbf{h}^{(i)}$$

- e.g., a good weight vector

The restaurant was expensive , but the menu was great

0.01 0.01 0.01 0.01 0.01 0.01 0.02 0.45 0.02 0.45



Sentiment Classification

- Attention-based LSTM
 - how to obtain the weight vector?
 - calculate the correlation between each word and the given aspect and then normalize

$$s_t = \mathbf{h}^{(t)\top} W \mathbf{h}_{food}$$

$$\alpha_i = \text{softmax}(a_t) = \frac{\exp(s_t)}{\sum_i \exp(s_i)}$$



Attention

- Formal definition: (K-Q-V).


- given query $\mathbf{q}$ and keys $\mathbf{x}_{1:N}$

$$\begin{aligned}\alpha_i &= p(z = i | \mathbf{x}_{1:N}, \mathbf{q}) \\ &= \text{softmax} \left(s(\mathbf{x}_i, \mathbf{q}) \right) \\ &= \frac{\exp \left(s(\mathbf{x}_i, \mathbf{q}) \right)}{\sum_{j=1}^N \exp \left(s(\mathbf{x}_j, \mathbf{q}) \right)},\end{aligned}$$

- $s(\mathbf{x}_i, \mathbf{q})$ is the score function

$$s(\mathbf{x}_i, \mathbf{q}) = \mathbf{x}_i^T \mathbf{q},$$

$$s(\mathbf{x}_i, \mathbf{q}) = \mathbf{x}_i^T W \mathbf{q},$$





Attention

- Formal definition: (K-Q-V).

- given value $\mathbf{v}_{1:N}$

$$a = \sum_i \alpha_i \mathbf{v}_i$$

- a is the output.

- Formal definition: (K-Q-V).


- given value $\mathbf{v}_{1:N}$

- a is the output.



Attention

- Formal definition: (K-Q-V).






Attention in RNN

- Global Attention
- Local Attention
- Bi-directional Attention
- Self Attention
- Key Value Attention



Self-Attention


- Self-Attention





Self-Attention


- Self-Attention





Today's class


- From ML to DL
- Neuron and ANN
- DL models
 - RNN 循环神经网络
 - CNN 卷积神经网络
 - Encoder-Decoder 编码器-解码器网络
 - Attention 注意力机制
 - Transformer 网络





Transformer

- Transformer
 - New framework
 - core component
 - Self-attention
 - Multi-head attention
 - other components
 - Position embedding 位置编码
 - Layer norm 层归一化
 - Residual connection 残差连接
 - FNN
- Fit to parallel training





Transformer

- Multi-head attention



1) This is our input sentence*

Thinking
Machines


2) We embed each word*




3) Split into 8 heads. We multiply X or R with weight matrices



...





4) Calculate attention using the resulting $Q/K/V$ matrices




...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



...



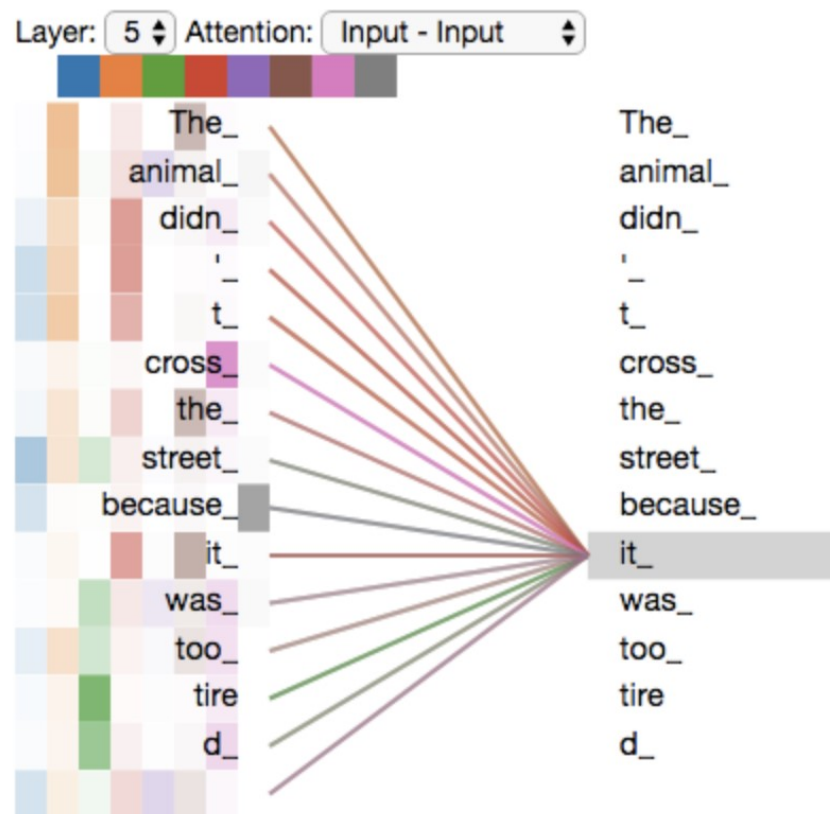

* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one





Transformer

- Multi-head attention





Transformer


- Position embedding
 - RNN consider the position information while self-attention not.






Transformer

- Layer norm
 - Increasing the convergence speed.
- Residual connection
 - Similar to ResNet, mitigate the Vanishing gradient problem.





Transformer



- Transformer apply masks to the input of first multi-head attention module to avoid seeing potential “future” sequence.
- Transformer may be applied to sequence transformation, and different input/output modal



Reference

1. 邱锡鹏. 神经网络与深度学习.
2. 自然语言处理中的自注意力模型 <https://xpqiu.github.io/slides/20200613-CAAI-%E8%87%AA%E7%84%B6%E8%AF%AD%E8%A8%80%E5%A4%84%E7%90%86%E4%B8%AD%E7%9A%84%E8%87%AA%E6%B3%A8%E6%84%8F%E5%8A%9B%E6%A8%A1%E5%9E%8B.pdf>
3. Recurrent Neural Networks and Language Models.
<http://web.stanford.edu/class/cs224n/slides/cs224n-2020-lecture06-rnnlm.pdf>
4. Vanishing Gradients and Fancy RNNs. <http://web.stanford.edu/class/cs224n/slides/cs224n-2020-lecture07-fancy-rnn.pdf>
5. Tixier A J P. Notes on deep learning for nlp[J]. arXiv preprint arXiv:1808.09772, 2018.



The next lecture

- Lecture 5: Language Model