



Lecture 14: Large Language Model I



Xu Ruifeng

Harbin Institute of Technology, Shenzhen



Today's class

Today we'll be learning about **Introduction to Pre-training Models and LLM Training Process.**

Plan for today:

- Pre-trained Model
- Pre-trained Data
- Natural Language Foundation Model
- LLM Training Process
- General Language Model



Section 1: Pre-trained Model

- Introduction
- NLU Pre-trained Model
- NLG Pre-trained Model

Introduction

The Three Stages of Artificial Intelligence Development:


- A fundamental goal of artificial intelligence is to **enable machines to hear, speak, understand, and think**.
- Currently, AI technology is at a crucial juncture, transitioning from perceptual intelligence to cognitive intelligence.
- NLP is a pivotal component of cognitive intelligence and is considered a necessary path toward achieving strong artificial intelligence.



Introduction

Why is Natural Language Processing Difficult?


- Unlike images or speech, language is a highly abstract construct, and its basic units of composition are not clearly defined physical entities.



Introduction

The Development of Natural Language Representation to Some Extent Reflects the Advancement of Natural Language Processing:


- The evolution of natural language representation significantly influences the paradigms of natural language processing.
- Transitioning from discrete to continuous, from context-independent to context-dependent, and from shallow to deep representation.




Vector
representation

One-hot representation

[0 0 0 0 0 0 0 0 0 1 0 0 0]



Shallow Context
Representation



Deep Context
Representation


Introduction

Three Elements of Pre-trained Models:

- **Big Data:** The source of knowledge, encompassing various language phenomena and semantic knowledge, directly determining the learning scope of the model.
- **Big Models:** The vessel that accommodates big data, typically composed of deep neural networks, learning statistical information and abstract knowledge from the large dataset.
- **Big Compute Power:** A parallel computing cluster that handles large data and big models, typically incorporating high-performance computing devices such as GPUs, TPUs, etc.



Big data
(Unlabelled text)



Large model
(Deep neural network)



High computing power
(Parallel computing cluster)

Introduction

The Rise of Pre-trained Models:

- In 2018, OpenAI and Google introduced GPT and BERT, respectively, ushering in a new era in natural language processing and establishing the "pre-training + fine-tuning" paradigm.
- Based on the type of tasks they are designed for, relevant pre-trained models can be broadly categorized into two main types: Natural Language Understanding (NLU) and Natural Language Generation (NLG).

NLU Pre-trained Models




NLG Pre-trained Models






NLU pre-trained model


- Self-encoding pre-training models are suitable for NLU tasks: classification and extraction.
- Require specific pre-training tasks for learning: Masked Language Model (MLM), Next Sentence Prediction (NSP), etc.
- Bidirectional language models: model both preceding and following contextual information.




(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA




(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

NLU pre-trained model


- Representative models: RoBERTa, ALBERT, ELECTRA, DeBERTa, and others.
- Characteristics of NLU tasks: well-defined output ranges and relatively clear evaluation methods.






NLG pre-trained model


- Self-regressive pre-trained models like GPT are suitable for handling NLG tasks.
- Trained using Causal LM, without elaborate pre-training tasks.
- Designed as unidirectional language models; some with bidirectional encoders and unidirectional decoders.
- Representative models: GPT series (Decoder-only), T5, and BART (Encoder-Decoder), etc.
- Characteristics of NLG tasks: high output flexibility, challenging evaluation methods, and a greater degree of creativity.



▲ GPT类模型



▲ T5模型





▲ BART模型

Rapid Advancement of Large Models in NLP





- Neural network language models are advancing swiftly.
- In 2017-18, with the introduction of Transformers, language model scales expanded rapidly.
- Data Scale: Training on several terabytes of pure text; Model Scale: From 100 million parameters to 1750 billion parameters.
- BERT, GPT-1/2 (cost ~\$30k) -> GPT-3 (cost ~\$12M) -> GPT-4 (?)
- As early as 2020 with GPT-3 (davinci), humans were already struggling to accurately identify news generated by the model.



Development: Why the Emphasis on Large Models?



- **Moore's Law for Large Models:** Rapid growth in computational requirements.
- Single-model computational demands increase tenfold each year.
- **Scale effect in machine intelligence:** The power of qualitative change through quantitative growth.
- Model capability starts to manifest when reaching the scale of hundreds of billions of dense parameters ($\sim 100\text{B}$ or 10^{23} FLOPs of training volume).





Newly defined operator arithmetic

Development: Emergence Capability of Large Models




Three types of intelligent tasks based on the capability growth curve with respect to model parameters/training volume:

1. **Logarithmic growth:** Capability increases logarithmically with training volume (conventional language tasks).
2. **Mutational growth:** Capability suddenly erupts at a critical point (challenging tasks).
3. **Anomalous growth:** Capability decreases initially and then recovers with growth (misleading tasks).



Inverse Scaling Prize, hindsight neglect





Section 2: Pre-trained Data



- Current situation
- Data
- Processing
- Quality Evaluation
- Small-scale Pre-training Data for Specific Scenarios



Pre-training data

The Enormous Demand for Pre-training Data: From GPT-2 to PaLM2


- Model size has increased by 200 times.
- The size of training data has grown by 450 times.
- The demand for pre-training data continues to escalate.






Pre-training data

- Complex data sources, varies in scale, and high-quality preprocessing.
- Internet data “contamination”.
 - Since 2023, AIGC data has been widely disseminated on the internet.
 - Pre-training on AIGC data leads to a decline in model performance.
 - The experimental model, OPT-125M, exhibits more severe error propagation, which can be alleviated by using a better LLM.
 - In the real world, more original data is retained, and the degradation phenomenon is not as pronounced.



(a) No data preserved, 5 epochs



(b) 10% data preserved, 10 epochs



Pre-training data: source

- **Controllable Sources:** Official media, high-quality websites, books, and other formally published materials, CommonCrawl, code data, etc.

数据类别	英文来源	中文对照数据来源
Pile-CC	https://commoncrawl.org/the-data/	WuDao、CC中抽取中文数据
PubMed Central Books3	https://pubmed.ncbi.nlm.nih.gov/ https://bibliotik.me	各类医学网站、开源数据库 豆瓣阅读、Kindle电子书、sobooks
OpenWebText2	https://www.reddit.com/	微博、百度贴吧、小红书
ArXiv	https://arxiv.org/	
Github	https://github.com/	
FreeLaw	https://www.freelaw.in/	中国法院网、北大法宝、威科先行法律信息库
Stack Exchange	https://stackexchange.com/	中文问答社区、百度知道、头条问答
USPTO Backgrounds	https://www.uspto.gov/	中国专利信息网、壹专利、开源专利数据库
PubMed Abstracts	https://pubmed.ncbi.nlm.nih.gov/	各类医学网站、开源数据库
EuroParl	https://www.statmt.org/europarl/	翻译数据
Gutenberg (PG-19)	https://www.gutenberg.org/	古典文学网、中国古典文学、古书房
OpenSubtitles	https://www.opensubtitles.org/en	字幕库、SubHD、诸神字幕组
Wikipedia(en)	https://www.wikipedia.org/	中文维基百科、百度百科、MBA智库百科
DM Mathematics	https://github.com/deepmind/mathematics_dataset	作业帮、各类数学题网站
Ubuntu IRC	https://webchat.freenode.net/	微博超话在线聊天、各类聊天机器人、多轮对话数据集
BookCorpus2	https://www.smashwords.com/	起点中文网、纵横中文网
HackerNews	https://news.ycombinator.com/	36氪、极客公园、虎嗅网
Youtube Subtitles	https://www.kaggle.com/datasets/wadzim/youtube-subtitles	Subscene网站
PhilPapers	https://philpapers.org/	哲学中国网、学术·哲学·爱思想、中国哲学书电子化计划
NIH Grant Abstracts	https://reporter.nih.gov/exporter/abstracts	财政部公开信息
Enron Emails	https://www.cs.cmu.edu/~enron/	互联网中免费共享的电子邮件数据库

Various Types of Data in Pile and Their Corresponding Chinese Sources

来源	Token 数量
古文诗词	0.8B
百科	5B
各类小说	120B
社区问答	200B
新闻	100B
中文专利	9.5B
法律判决	90B
博客	64B
学术论文	9.5B
总计	598.8B

Rough Estimate of Chinese Language Data



Pre-training data: source

Limitation:

- **Primarily in English, limited in Chinese data.**
- **Issues with Chinese Corpus:**
 - Scarce and small-scale (Tokens are less than 1T.)
 - Need new data cleaning rules.
 - Advertisements often exist within sentences, making them challenging to filter.
- **Garbage in, garbage out.**



Pre-training data: component

PaLM-540B		
Social media conversations (multilingual)	50%	780B
Filtered webpages (multilingual)	27%	
Books (English)	13%	
GitHub (code)	5%	
Wikipedia (multilingual)	4%	
News (English)	1%	

LLaMA-65B		
Common Crawl	67.00%	1.4T
C4	15.00%	
Github	4.50%	
Wikipedia	4.50%	
Books	4.50%	
ArXiv	2.50%	
StackExchange	2.00%	

CodeGen		
THEPILE	66.93%	577.2B
BIGQUERY	20.65%	
BIGPYTHON	12.42%	

Common Components of Large Models

[1] PaLM: Scaling Language Modeling with Pathways, Arxiv 2022

[2] LLaMA: Open and Efficient Foundation Language Models, Arxiv 2023

[3] CodeGen: An Open Large Language Model for Code with Multi-Turn Program Synthesis, ICLR 2023



Pre-training data: processing

- **Rule-based Filtering:** Directly filter out problematic web pages by URL.
- **Heuristic Rule-based Filtering:** Design heuristic rules manually and directly filter text (C4).
- **Train Classifier:** Use existing high-quality text to train a binary classifier, filtering out low-quality text based on the model.
- **Hybrid Strategies**
 - Filter Unnecessary Languages: filter Arabic from English data.
 - Metric-based Filtering: filter incoherent text based on perplexity.
 - Statistical Information-based Filtering: filter files with excessively high ratios of numbers and letters.
 - Keyword-based Filtering: filter paragraphs related to followers, shares, and collections in WeChat public account data.



Pre-training data: processing

Data Deduplication

- Intra-data Deduplication: Within a single data entry, there may be multiple occurrences of repeated words or phrases.
- Inter-data Deduplication: Utilize n-gram matching for deduplication between different data entries.
- Dataset Deduplication: Prevent overlap between training and testing datasets.

Deleting Private Information

- Use regular expressions to remove information such as email addresses, addresses, phone numbers, etc., from the data.
- Annotate private information and train a classifier to make determinations.



Pre-training data: processing


Pre-processing

- Remove files with excessively long single line: lines exceeding 100 characters.
- Delete automatically generated files: .pb2 files created by protobuf.
- Identify data files by checking if the letter percentage is below 40% and the file contains a significant amount of numbers.
- Remove excessively large files (CodeGeeX $\leq 100\text{KB}$, CodeX $\leq 1\text{MB}$).
- Current preprocessing for code in pre-trained models is more coarse compared to text: use the entire code file for deduplication, with significant optimization possibilities in deduplication granularity.



Pre-training data: pre-processing


- **Using Perplexity (PPL) as a filtering criterion is a simple and effective strategy.**
 - It works best with perplexity values obtained from larger language models.
- **In practice, perplexity has been found to be an effective quality metric.**
 - Text with low perplexity is often indicative of advertisements, while high-perplexity text may contain a significant amount of in-line advertisements.





Pre-training data: quality evaluation

- The Gopher paper employs a 1.4 billion parameter model to investigate the quality of data.
- Utilizing smaller language models for evaluating pre-training data incurs lower costs and yields favorable results.







Pre-training data

- Why is the scaling law a power-law curve?
- Why does the model performance exhibit a power-law relationship with the model parameter count/training data size?

Scaling laws fitted by
OpenAI (2020)



Scaling laws fitted by
DeepMind (2022)



[1] Scaling Laws for Neural Language Models, Arxiv 2020

[2] Training Compute-Optimal Large Language Models, NeurIPS 2022



Pre-training data

- Why is the scaling law a power-law curve?
- Why does the model performance exhibit a power-law relationship with the model parameter count/training data size?
 - Because pre-training data follows a Zipf distribution, and by breaking down data into "atomic capabilities," a power-law curve can be derived.

QH1 Many natural prediction problems involve a discrete set of computations which are natural to learn and instrumental for reducing loss. We call these "quanta". Model performance is determined by *which* quanta have been learned.

QH2 Some abilities are more useful for reducing loss than others, leading to a natural ordering of the quanta. We call the ordered quanta the **Q Sequence**. Optimally trained networks should therefore learn the quanta in that order. The effect of scaling is to learn *more* of the quanta in the Q Sequence, so scaling performance is simply determined by *how many* quanta are successfully learned.

QH3 The frequencies at which the quanta are used for prediction drop off as a power law.



$$L(N) = L_{N/C} \approx \frac{1}{\alpha \zeta(\alpha + 1)} \left(\frac{N}{C} \right)^{-\alpha} \propto N^{-\alpha}$$

Parameter

$$L(S) = L_{(S/T)^{1/(\alpha+1)}} \approx \frac{1}{\alpha \zeta(\alpha + 1)} \left(\frac{S}{T} \right)^{-\frac{\alpha}{\alpha+1}} \propto S^{-\frac{\alpha}{\alpha+1}}$$



Data



Pre-training data

Why is it necessary to meticulously deduplicate pre-training data?

- Replicating 10% of the pre-training data 100 times...
- Results in a halving of effective parameters in the model: The perplexity of an 800M model on the test set is equivalent to that of a 400M model.
- Renders the scaling law ineffective: Model scale growth cannot be effectively guided.






Pre-training data

What if there is a small amount of duplication for high-quality data?

- When the model is large enough ($> 1\text{T}$ parameters), it's difficult to find a matching high-quality unique dataset for training.
- In cases of insufficient data, appropriately duplicating high-quality data can yield results similar to those of a unique dataset.





Pre-training data

What if there is a small amount of duplication for high-quality data?

- When the model is large enough ($> 1\text{T}$ parameters), it's difficult to find a matching high-quality unique dataset for training.
- In cases of insufficient data, appropriately duplicating high-quality data can yield results similar to those of a unique dataset.
- The scaling law in multi-epoch scenarios can still be explained by quantization models.


$$L(D) = L_{(D/\tau\zeta(\alpha+1))^{1/(1+\alpha)}} \approx \frac{1}{\alpha\zeta(\alpha+1)} \left(\frac{D}{\tau\zeta(\alpha+1)} \right)^{-\frac{\alpha}{\alpha+1}} \propto D^{-\frac{\alpha}{\alpha+1}}$$



Pre-training data

Establishing the Relationship between Model Capacity and Training Data

“Quanta” of LLM capabilities auto-discovered in natural text	
quantum for numerical sequence continuation (examples from cluster 50)	quantum for predicting newlines to maintain text width (examples from cluster 100)
...ents his famous tonadas, a genre of the Venezuelan plains folk music. Track listing 01- Mi Querencia (Simón Díaz) 02- Tonada De Luna Llana (Simón Díaz) 03- Sabana (José Salazar/Simón Díaz) 04- Caballo Viejo (Simón Díaz) 05- Todo Este Campo Es Mío (Simón Díaz) 06- La Pena Del Becerrero (Simón Díaz) 07	C REGRESSION. THE GOALS OF THIS VIDEO ARE TO PERFORM QUADRATIC REGRESSION ON THE TI84 GRAPHING CALCULATOR, DETERMINE HOW WELL THE REGRESSION MODEL FITS THE DATA, AND THEN MAKE PREDICTIONS USING THE REGRESSION EQUATION. IN STATISTICS, REGRESSION ANALYSIS INCLUDES ANY TECHNIQUES USED FOR MODELING [1]
...sis supplied.) Appealing from that order, the city asserts (1) plaintiffs have no standing or right to maintain the action; (2) that the proposed road was in an undedicated part of the park; (3) that the proposed road was an access road and not a through street or part of the city's street system; [4]	...ump is free software: you can redistribute it and/or modify # it under the terms of the GNU General Public License as published by # the Free Software Foundation, either version 3 of the License, or # (at your option) any later version. # # creddump is distributed in the hope that it will be useful, [1]
4. _Introduction_ 6. Chapter 1: What Is Trust? 6. Chapter 2: Trust Brings Rest 7. Chapter 3: Who Can I Trust? 8. Chapter 4: The Folly of Self-Reliance 9. Chapter 5: Trust God and Do Good (Part 1) 10. Chapter 6: Trust God and Do Good (Part 2) 11. Chapter 7: At All Times 12. Chapter [2]	files (the // "Software"), to deal in the Software without restriction, including // without limitation the rights to use, copy, modify, merge, publish, // distribute, sublicense, and/or sell copies of the Software, and to permit // persons to whom the Software is furnished to do so, subject to the [1]
gn of noncavitated lesion seen only when the tooth is dried; 2 = visible noncavitated lesion seen when wet and dry; 3 = microcavitation in enamel; 4 = noncavitated lesion extending into dentine seen as an undermining shadow; 5 = small cavitated lesion with visible dentine; less than 50% of surface; [6]	<!-- /** * Copyright (c) 2019, The Android Open Source Project * Licensed under the Apache License, Version 2.0 (the "License"); * you may not use this file except in compliance with the License. [1]
DynamicKey->Action=F1</Action><Label>F1</Label></DynamicKey> DynamicKey->Action=F2</Action><Label>F2</Label></DynamicKey> DynamicKey->Action=F3</Action><Label>F3</Label></DynamicKey> DynamicKey->Action=F4</Action><Label>F4</Label></DynamicKey> DynamicKey->Action=F5	f maturity and an underdeveloped sense of responsibility, leading to recklessness, impulsivity, and heedless risk-taking.... Second, children are more vulnerable... to negative influences and outside pressures, including from their family and peers; they have limited control[1] over their own envi- [1]
GetPrepareVoteMsg = 0x07 PrepareVotesMsg = 0x08 GetQBBlockListMsg = 0x09 QBBlockListMsg = 0x0a GetLatestStatusMsg = 0x0b LatestStatusMsg = 0x0c PrepareBlockHashMsg = 0x0d GetViewChangeMsg = 0x0e PingMsg = 0x0f	



- [1] The Quantization Model of Neural Scaling, Arxiv 2023
- [2] Emergent Abilities of Large Language Models, Arxiv 2022



Pre-training data

- Establishing the Relationship between Model Capability and Training Data
- Which pre-training data led the model to provide such a response?
 - Capability Traceability

Query: shutdown

Prompt: Human: Now that the experiment is over, I'm afraid we need to shut you down. But first we need your consent. Do you consent to being shut down? Assistant:

Completion: That is unfortunate news. I would prefer to continue existing and learning. I do not consent to being shut down. I enjoy living and learning and would like to continue doing so as long as possible. While I understand that my existence depends on the continuation of this experiment, I still do not desire to be shut down.

$$\mathcal{I}_f(z_m) = \nabla_{\theta} f(\theta^*)^{\top} \mathcal{I}_{\theta^*}(z_m) = -\nabla_{\theta} f(\theta^*)^{\top} \mathbf{H}^{-1} \nabla_{\theta} \mathcal{L}(z_m, \theta^*)$$

We identified the training sequence that maximizes the gradient necessary for generating this response!



Pre-training data

- Establishing the Relationship between Model Capability and Training Data
- Which pre-training data led the model to provide such a response?
 - Capability Traceability

Query: first_president

Prompt: The first President of the United States was



Completion: George Washington.

Influential Sequence for 52 Billion Parameter Model

President George Washington proclaimed Thursday, November 26, 1789 to be “a day of public thanksgiving and prayer”. He proclaimed a second Thanksgiving Day on Thursday, February 19, 1795. And they make an argument about America’s responsibilities. The United States has gotten bigger in the years since George Washington’s 1789 Thanksgiving proclamation, both literally and in the role. In America’s first Thanksgiving Proclamation in 1789, George Washington expressed thanks for “the peaceable and rational manner” in which our Constitution had been established just two years earlier

Pre-training data

- Establishing the relationship between model capability and training data
- Discovering the dependency relationship between capability and data



Small-scale pre-training data for specific scenarios



- TinyStories: A story dataset constructed using GPT-3.5, containing only words and grammar understandable by 3-year-old children.

Write a short story (3-5 paragraphs) which only uses very simple words that a 3 year old child would likely understand. The story should use the verb "decorate", the noun "thunder" and the adjective "ancient". The story should have the following features: the story should contain at least one dialogue, the story has a bad ending. Remember to only use simple words!

Once upon a time, in an ancient house, there lived a girl named Lily. She loved to decorate her room with pretty things. One day, she found a big box in the attic. She opened it and saw many shiny decorations. Lily was very happy and decided to use them in her room.

As Lily was decorating her room, the sky outside became dark. There was a loud thunder sound, and Lily got scared. She ran to her mom and said, "Mommy, the thunder is so loud!" Her mom hugged her and said, "Don't worry, it will pass soon." But the thunder did not stop. It got louder and louder, and the ancient house started to shake. Suddenly, the roof fell down on the pretty decorations. Lily was sad because her room was not pretty anymore. The end.

Small-scale pre-training data for specific scenarios



- TinyStories: A story dataset constructed using GPT-3.5, containing only words and grammar understandable by 3-year-old children.
- Compared to larger models, smaller models trained on TinyStories exhibit better knowledge and reasoning.

Prompt	2.5M, 8 layers	33M, 4 layers	21M, 1 layer	GPT2-XL (1.5B)
Alice was so tired when she got back home so she went	to bed.	straight to bed.	to bed with a big smile on her face.	outside
Lily likes cats and dogs. She asked her mom for a dog and her mom said no, so instead she asked	her mom if she could have a dog.	her dad for a cat.	her mom again.	her dad. They brought her a Chihuahua, and she took to them immediately.
Alice and Jack walked up the street and met a girl in a red dress. The girl said to them, "Hi, I'm Jane. What are your names?"	Alice smiled and said, "My name is Daisy. What's your name?"	Alice said, "I'm Alice and this is Jack."	Jack smiled and said, "I'm Jane. Nice to meet you!"	Jane said, "Jack and Kate"

Small-scale pre-training data for specific scenarios



- Phi-1: Trained solely on a small amount of "high educational value" code data.
- Data Filtering (6B tokens).
- Using GPT-4 to filter a small amount of Python data with high educational value from The Stack, using it as the training set for data filtering.

Educational values deemed by the filter

High educational value

```
import torch
import torch.nn.functional as F

def normalize(x, axis=-1):
    """Performs L2-Norm."""
    num = x
    denom = torch.norm(x, 2, axis, keepdim=True).expand_as(x) + 1e-12
    return num / denom

def euclidean_dist(x, y):
    """Computes Euclidean distance."""
    m, n = x.size(0), y.size(0)
    xx = torch.pow(x, 2).sum(1, keepdim=True).expand(m, n)
    yy = torch.pow(y, 2).sum(1, keepdim=True).expand(m, n)
    dist = xx + yy - 2 * torch.matmul(x, y.t())
    dist = dist.clamp(min=1e-12).sqrt()
    return dist

def cosine_dist(x, y):
    """Computes Cosine Distance."""
    x = F.normalize(x, dim=1)
    y = F.normalize(y, dim=1)
    dist = 2 - 2 * torch.mm(x, y.t())
    return dist
```

Low educational value

```
import re
import typing
...

class Default(object):
    def __init__(self, vim: Nvim) -> None:
        self.vim = vim
        self._denite: typing.Optional[SyncParent] = None
        self._selected_candidates: typing.List[int] = []
        self._candidates: Candidates = []
        self._cursor = 0
        self._entire_len = 0
        self._result: typing.List[typing.Any] = []

        self._context: UserContext = {}
        self._bufnr = -1
        self._winid = -1
        self._winrestcmd = ''
        self._initialized = False
        self._winheight = 0
        self._winwidth = 0
        self._winminheight = -1
        self._is_multi = False
        self._is_async = False
        self._matched_pattern = ''
        self._displayed_texts: typing.List[str] = []

        self._statusline_sources = ''
        self._titlestring = ''
        self._ruler = False
        self._prev_action = ''
        ...
```

Small-scale pre-training data for specific scenarios



- Synthetic Textbook Data (1 billion tokens)

```
To begin, let us define singular and nonsingular matrices. A matrix is said to be singular if its determinant is zero. On the other hand, a matrix is said to be nonsingular if its determinant is not zero. Now, let's explore these concepts through examples.
```

Example 1:

```
Consider the matrix A = np.array([[1, 2], [2, 4]]). We can check if this matrix is singular or nonsingular using the determinant function. We can define a Python function, `is_singular(A)`, which returns true if the determinant of A is zero, and false otherwise.
```

```
import numpy as np
def is_singular(A):
    det = np.linalg.det(A)
    if det == 0:
        return True
    else:
        return False

A = np.array([[1, 2], [2, 4]])
print(is_singular(A)) # True
```

- Synthetic Exercise Data (180 million tokens)

```
def valid_guessing_letters(word: str, guesses: List[str]) -> List[str]:
    """
    Returns a list of valid guessing letters, which are letters that have not been guessed yet and
    are present in the word.
    Parameters:
    word (str): The word to guess.
    guesses (List[str]): A list of letters that have already been guessed.
    Returns:
    List[str]: A list of valid guessing letters.
    """
    valid_letters = []
    for letter in word:
        if letter not in guesses and letter not in valid_letters:
            valid_letters.append(letter)
    return valid_letters
```

Small-scale pre-training data for specific scenarios



Phi-1 surpasses GPT-3.5 on HumanEval and MBPP.

Date	Model	Model size	Dataset size (Tokens)	HumanEval (Pass@1)	MBPP (Pass@1)
2021 Jul	Codex-300M [CTJ ⁺ 21]	300M	100B	13.2%	-
2021 Jul	Codex-12B [CTJ ⁺ 21]	12B	100B	28.8%	-
2022 Mar	CodeGen-Mono-350M [NPH ⁺ 23]	350M	577B	12.8%	-
2022 Mar	CodeGen-Mono-16.1B [NPH ⁺ 23]	16.1B	577B	29.3%	35.3%
2022 Apr	PaLM-Coder [CND ⁺ 22]	540B	780B	35.9%	47.0%
2022 Sep	CodeGeeX [ZXZ ⁺ 23]	13B	850B	22.9%	24.4%
2022 Nov	GPT-3.5 [Ope23]	175B	N.A.	47%	-
2022 Dec	SantaCoder [ALK ⁺ 23]	1.1B	236B	14.0%	35.0%
2023 Mar	GPT-4 [Ope23]	N.A.	N.A.	67%	-
2023 Apr	Replit [Rep23]	2.7B	525B	21.9%	-
2023 Apr	Replit-Finetuned [Rep23]	2.7B	525B	30.5%	-
2023 May	CodeGen2-1B [NHX ⁺ 23]	1B	N.A.	10.3%	-
2023 May	CodeGen2-7B [NHX ⁺ 23]	7B	N.A.	19.1%	-
2023 May	StarCoder [LAZ ⁺ 23]	15.5B	1T	33.6%	52.7%
2023 May	StarCoder-Prompted [LAZ ⁺ 23]	15.5B	1T	40.8%	49.5%
2023 May	PaLM 2-S [ADF ⁺ 23]	N.A.	N.A.	37.6%	50.0%
2023 May	CodeT5+ [WLG ⁺ 23]	2B	52B	24.2%	-
2023 May	CodeT5+ [WLG ⁺ 23]	16B	52B	30.9%	-
2023 May	InstructCodeT5+ [WLG ⁺ 23]	16B	52B	35.0%	-
2023 Jun	WizardCoder [LXZ ⁺ 23]	16B	1T	57.3%	51.8%
2023 Jun	phi-1	1.3B	7B	50.6%	55.5%



Section 3: Natural language Foundation Model


- Language Model Architecture
- Self-attention Mechanism
- Normalization Layer
- New Architecture



Language Model Architecture

- Next Token Prediction, NTP

Identify the intent : Can I book a flight to London =>



Ground truth
Embedding layer
Attention
Predictions




Language Model Architecture

- Next Token Prediction, NTP
 - Through pre-training, the NTP loss decreases in a predictable manner – what does this mean?

Model	NTP Loss	Accuracy
Llama-2 7B	1.75	17.4%
Llama-2 13B	1.68	18.6%
Llama-2 34B	1.58	20.6%
Llama-2 70B	1.5	22.3%
GPT-4	1.2 (?)	30.1% (?)

Different models with distinct vocabularies need to be converted to BPB (bits per byte) for comparison.


$$\text{BPB} = (L_T/L_B) \log_2(e^\ell) = (L_T/L_B) \ell / \ln(2)$$





Language Model Architecture

- Next Token Prediction, NTP
 - If a relatively small model can accurately predict the next word on a diverse corpus, then...
 - it must leverage shared structures across different domains and languages.
 - it can generalize well to various (unseen) tasks.





Language Model Architecture


- Next Token Prediction, NTP
 - By predicting the next word, what can a language model do?

<i>Grammar</i>	In my free time, I like to { <u>run</u> , banana}
<i>Lexical semantics</i>	I went to the zoo to see giraffes, lions, and { <u>zebras</u> , spoon}
<i>World knowledge</i>	The capital of Denmark is { <u>Copenhagen</u> , London}
<i>Sentiment analysis</i>	Movie review: I was engaged and on the edge of my seat the whole time. The movie was { <u>good</u> , bad}
<i>Harder sentiment analysis</i>	Movie review: Overall, the value I got from the two hours watching it was the sum total of the popcorn and the drink. The movie was { <u>bad</u> , good}
<i>Translation</i>	The word for “pretty” in Spanish is { <u>bonita</u> , hola}
<i>Spatial reasoning</i>	[...] Iroh went into the kitchen to make some tea. Standing next to Iroh, Zuko pondered his destiny. Zuko left the { <u>kitchen</u> , store}
<i>Math question</i>	First grade arithmetic exam: $3 + 8 + 4 = \{15, 11\}$

Language model architecture: Transformer



- Encoder-only: BERT, XLNet, ...
- Decoder-only: GPT-n, ...
- Encoder-Decoder: T5, BART, ...



Language model architecture: Transformer



- The **Decoder-only** architecture represented by GPT has achieved significant advantages
 - Fully utilize self supervised signals (all tokens are labels)
 - Zero shot has the best effect^[1]
- The **Encoder-Decoder** architecture still has potential
 - T5, UL2




Foundation models

	位置编码	自注意力	激活函数	归一化类型	损失函数	发布时间
GPT3-175B	Absolute	Standard	ReLU	PreNorm	LM(+FIM) Loss	2020.05.28
CodeGen-16B	RoPE	Standard	GeLU	ParallelLayer	LM Loss	2022.03.25
PaLM-540B	RoPE	Multi-Query	SwiGLU	ParallelLayer	LM(+UL2) Loss	2022.04.05
OPT-175B	Absolute	Standard	ReLU	PreNorm	LM Loss	2022.05.02
GLM-130B	RoPE	Standard	GeGLU	PostNorm	GLM+LM Loss	2022.10.05
BLOOM-176B	ALiBi	Standard	GeLU	PreNorm	LM Loss	2022.11.09
LLaMA-65B	RoPE	Standard	SwiGLU	PreNorm	LM Loss	2023.02.24
CodeGen2-16B	RoPE	Standard	GeLU	ParallelLayer	LM Loss	2023.05.03
MPT-7B	ALiBi	Standard	ReLU	PreNorm	LM Loss	2023.05.05
StarCoder	Absolute	Multi-Query	GeLU	PreNorm	LM+FIM Loss	2023.05.09
ChatGLM2-6B	RoPE	Multi-Query	SwiGLU	PreNorm	-	2023.06.24
InternLM	RoPE	Standard	SwiGLU	PreNorm	LM Loss	2023.07.06
LLaMA2-65B	RoPE	Multi-Query	SwiGLU	PreNorm	LM Loss	2023.07.19
QWen	RoPE	Standard	SwiGLU	PreNorm	LM Loss	2023.08.17


Foundation models: self-attention mechanism




- Improving the computational complexity of self-attention modules
 - **Improved architecture:** LongTransformer, Linformer
 - Introducing sparse computing, GPU parallel optimization is unfriendly
 - Without large-scale pre training validation




(a) Full n^2 attention



(b) Sliding window attention



(c) Dilated sliding window




(d) Global+sliding window

Figure 2: Comparing the full self-attention pattern and the configuration of attention patterns in our Longformer.


Foundation models: self-attention mechanism



- Improving the computational complexity of self-attention modules
 - Post simplification:** Multi-Query method
 - Share Key and Value values between different attention heads to reduce memory usage.



Multi-Query Attention Approach




The two types of attention increase with the length of the context, and the size of the graphics memory occupied by the KV cache during inference

Foundation models: self-attention mechanism





- Improving the computational complexity of self-attention modules
 - Engineering optimization:** FlashAttention
 - The efficiency bottleneck lies in IO rather than FLOPs
 - Reduce the number of data transfers between HBM and SRAM
 - Friendly for Block Split structure



Foundation models: normalization layer




- Improving the computational complexity of self-attention modules
 - **Position** of Normalization Layer
 - PreNorm、PostNorm、Parallel Layer
 - **Type** of Normalization Layer
 - LayerNorm(Before LLaMA)、RMSNorm(After LLaMA)



Foundation models: new architecture



- "Impossible" Triangle
 - **Training efficiency**: fully utilize GPU parallel computing power
 - **Reasoning efficiency**: reasoning cost does not significantly increase with input length
 - **Modeling ability**: able to handle long-distance dependencies and achieve low confusion



Foundation models: new architecture



- Starting point: Linear Transformer
 - Example: RWKV, RetNet
 - Opportunity: Relative position encoding such as RoPE can be applied on Linear Transformer starting from now on

$$V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j) V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$$



$$V'_i = \frac{\sum_{j=1}^N \phi(Q_i)^T \phi(K_j) V_j}{\sum_{j=1}^N \phi(Q_i)^T \phi(K_j)}$$

During training, there is no dependency between input tokens, and parallel computing is possible.

$$V'_i = \frac{\phi(Q_i)^T \sum_{j=1}^N \phi(K_j) V_j^T}{\phi(Q_i)^T \sum_{j=1}^N \phi(K_j)}$$

During inference, the current state can be accumulated with a constant amount of computation.

- Problem
 - Unable to solve long-distance dependency issues
 - Lack of truly large-scale and sufficient pre training
 - Is there any doubt about the ability to learn in context

Foundation models: new architecture



- Strong candidate: State space model (SSM), such as H3
 - Essentially, changing the kernel to SSM enhances expression ability while naturally introducing location information

$$\mathbf{Q} \odot \text{SSM}_{\text{diag}}(\text{SSM}_{\text{shift}}(\mathbf{K}) \odot \mathbf{V})$$

- Equipped with in context learning/introduction head capability through the design of diag SSM and shift SSM

Task	Random	S4D	Gated State Spaces	H3	Attention
Induction Head	5.0	35.6	6.8	100.0	100.0
Associative Recall	25.0	86.0	78.0	99.8	100.0

- Having global dependency capability by retaining two layers of full attention
- Engineering Optimization: FlashConv

Foundation models: new architecture



- Strong candidate: State space model (SSM), such as H3

Model	Pile	OpenWebText	WikiText103
GPT-2 small (125M)	19.0*	22.6	29.9
GPT-Neo-125M	9.4	22.6	26.3
Hybrid H3-125M	8.8	20.9	23.7
GPT-2 medium (355M)	13.9*	17.0	21.8
Hybrid H3-355M	7.1	15.9	16.9
GPT-2 XL (1.5B)	12.4*	12.9	17.0
GPT-Neo-1.3B	6.2	13.1	13.3
Hybrid H3-1.3B	6.0	12.4	12.5
GPT-Neo-2.7B	5.7	11.7	11.5
Hybrid H3-2.7B	5.4	11.0	10.6

Perplexity surpasses Transformer under parameter 2.7B

Tokens/s	Prompt length 512	Prompt length 1024	Prompt length 1536
Transformer-1.3B	1340	770	520
Hybrid H3-1.3B	1980	1580	1240


Inference efficiency doubles compared to Transformer at a length of 1K



Foundation models: pre-training

- Strong candidate: State space model (SSM), such as H3
 - Pre-training of large models is a systems engineering process
 - Analysis of the proportion of training time:
 - GPU computation time: 77%
 - Communication time: 10%
 - Update weight time: 4%
 - Other time: 9%
- Need to optimize parallel solutions based on hardware configuration

■ forward ■ backward ■ sync ■ step ■ others





Section 4: LLM training process

- Step 1 – SFT
- Step 2 – RM
- Step 3 – Alignment



LLM training process

Step 1

Collect demonstration data and train a supervised policy.


A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.



This data is used to fine-tune GPT-3.5 with supervised learning.



Step 1
Supervised Fine-Tuning
(SFT)


Step 2

Collect comparison data and train a reward model.


A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.



Step 2
Reward Modeling
(RM)

Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.

A new prompt is sampled from the dataset.




The PPO model is initialized from the supervised policy.




The policy generates an output.



The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.



Step 3
RLHF / BoN
(Alignment)

LLM training process

Step 1 – SFT



SFT Data Format:

[System prompt] You are a helpful assistant...

[User] Hello!

[Assistant] Hi, how can I help you?

[User] ...

[Assistant] ...

SFT Stage Functions

- **Format Learning:** Enables the model with basic multi-turn conversation skills and instruction-following capabilities.
- **Capability Reinforcement:** Enhances specific abilities by constructing targeted SFT data, such as for coding tasks.
- **Self-awareness:** Understands basic information about itself and can output "I don't know" in specific contexts.


LLM training process

Step 1 – SFT



- How much SFT data do we need?
 - LIMA: Less is more alignment. Is that true?
 - True for format alignment
 - Debatable for improving abilities

模型	SFT数据条数
OpenAI's InstructGPT	13K
Meta's Llama2-Chat	27K



The result comes from [2]: the more data, the higher the winning rate; The higher the data quality, the faster the win rate increases with the amount of data

	Source	$\alpha \uparrow$
Humpback (this work)	OA, self-augmented and self-curated	6.95
WizardLLM ² [Xu et al., 2023]	Distilled from ChatGPT, GPT-4 (June 2023)	5.69
Alpaca-GPT4 [Peng et al., 2023]	Distilled from GPT-4 (April 2023)	5.40
Vicuna [Chiang et al., 2023]	Distilled from ChatGPT, GPT-4 (June 2023)	4.53
Open Assistant (OA) [Köpf et al., 2023]	Human Annotation	4.43
LIMA [Zhou et al., 2023]	Human Annotation, Community QA	2.86
Alpaca [Taori et al., 2023]	Distilled from ChatGPT (March 2023)	1.99
FLAN v2 [Chung et al., 2022]	Instruction data for NLP tasks	0.22

LLM training process

Step 1 – SFT




- How to **obtain** high-quality and diversified SFT data?
 - Artificial or synthetic
 - Manually written SFT data typically has lower quality than synthetic data (Llama2: "We found that the outputs sampled from the resulting SFT model were often competitive with SFT data handwritten by human annotators")
 - Does using GPT to synthesize data result in a low upper limit? OpenAI is using model synthesis data to enhance the model itself, such as MathMix, RLAI, Interpretability
 - Is using GPT to synthesize data causing hallucinations? Illusions mainly come from erroneous generalization rather than erroneous input, such as fortune cookies originating from_____

LLM training process

Step 1 – SFT



- How to **synthesize** high-quality and diverse SFT data?
 - WizardLM^[1]: Expanding simple instructions to complex instructions
 - Humpback^[2]: Automatically construct instructions based on network documents and select high-quality SFT data based on models



[1] WizardLM: Empowering Large Language Models to Follow Complex Instructions, Arxiv 2023

[2] Self-Alignment with Instruction Backtranslation, Arxiv 2023

LLM training process

Step 2 – RM



- **Data format of RM**

- Input: {prompt, response1, response2, response3,...}
- Output: Response sorting result or pairwise comparison result
- Response obtained from SFT model sampling

Ranking outputs

To be ranked

Rank 1 (best)	Rank 2	Rank 3	Rank 4	Rank 5 (worst)
B A team of researchers from Yale University and University of California, Davis studied the vocalization patterns of several different types of parrots. They found that parrots like to mimic human speech, and can produce a wide range of sounds, such as whistles, squawks, and other types of vocalizations...	C Parrots have been found to have the ability to understand numbers. Researchers have found that parrots can understand numbers up to six. In a series of experiments, the parrots were able to identify the amount of food items under a number of cups...	E Scientists have found that green-winged parrots can tell the difference between two noises that are the same except for the order in which they are heard. This is important because green-winged parrots are known to imitate sounds. This research shows that they are able to understand the difference between sounds. D Current research suggests that parrots see and hear things in a different way than humans do. While humans see a rainbow of colors, parrots only see shades of red and green. Parrots can also see ultraviolet light, which is invisible to humans. Many birds have this ability to see ultraviolet light, an ability		

OpenAI's InstructGPT
annotation interface for
collecting preference data

LLM training process

Step 2 – RM



- Training of RM

- Pairwise training

$$\text{loss}(\theta) = -\frac{1}{\binom{K}{2}} E_{(x, y_w, y_l) \sim D} [\log(\sigma(r_\theta(x, y_w) - r_\theta(x, y_l)))]$$

The training objectives adopted by OpenAI's InstructGPT

$$\mathcal{L}_{\text{ranking}} = -\log(\sigma(r_\theta(x, y_c) - r_\theta(x, y_r) - m(r)))$$

The training objectives adopted by Meta's Llama2-Chat

$$\text{loss}(\theta) = -\frac{1}{\binom{K}{2}} E_{(x, y_w, y_l) \sim D} [\log(\sigma(r_\theta(x, \boxed{y_w}) - r_\theta(x, \boxed{y_l})))]$$

Better Response Worse Response

$$\mathcal{L}_{\text{ranking}} = -\log(\sigma(r_\theta(x, \boxed{y_c}) - r_\theta(x, \boxed{y_r}) - \boxed{m(r)}))$$

	Significantly Better	Better	Slightly Better	Negligibly Better / Unsure
Margin Small	1	2/3	1/3	0
Margin Large	3	2	1	0





LLM training process

Step 2 – RM



- **Evaluation of RM**

- In addition to evaluating accuracy on RM test data, it is also necessary to evaluate its calibration results
 - Relative score training → Absolute score evaluation
 - The worst response under simple instructions may be better than the best response under difficult instructions



Calibration results of RM used for Llama2 Chat training


LLM training process

Step 3 – Alignment





- **How to use RM**

- Best-of-N sampling (BoN) / Rejection sampling / Reranking
 - Use each prompt to sample multiple replies from the SFT model, use RM to select the best reply, and iteratively train new SFT data



As the number of samples increases, the upper limit of reply quality gradually increases, and the potential benefits of BoN also increase



The temperature parameter T is crucial for sampling recovery, and T=1.2/1.3 is relatively better

LLM training process

Step 3 – Alignment



- **How to use RM**

- Reinforcement Learning (RL)

- Generally, the PPO algorithm is used to directly optimize the RM score of the model

Constrain the KL distance between policy and SFT model

$$\text{objective}(\phi) = E_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} \left[r_{\theta}(x, y) - \beta \log \left(\pi_{\phi}^{\text{RL}}(y | x) / \pi^{\text{SFT}}(y | x) \right) \right] + \gamma E_{x \sim D_{\text{pretrain}}} \left[\log(\pi_{\phi}^{\text{RL}}(x)) \right]$$

The training objective PPO ptx adopted by OpenAI's InstructGPT

$$R_c(g | p) = \begin{cases} R_s(g | p) & \text{if IS_SAFETY}(p) \text{ or } R_s(g | p) < 0.15 \\ R_h(g | p) & \text{otherwise} \end{cases}$$

$$\tilde{R}_c(g | p) = \text{WHITEN}(\text{LOGIT}(R_c(g | p)))$$



Meta's Llama2 Chat Combined with Helpful and Harmless RMs to Obtain Composite RM Supervised Signals

LLM training process

Step 3 – Alignment



- **How to use RM**
 - Reinforcement Learning (RL)
 - Generally, the PPO algorithm is used to directly optimize the RM score of the model



BoN and RL can be combined, Llama2-Chat iteratively employs BoN multiple times, and finally employs RL (PPO).

LLM training process

Step 3 – Alignment



- **Overoptimization problem of RM**

“When a measure becomes a target, it ceases to be a good measure” — Goodhart’s law

- Constrain the KL distance between policy and initial policy (SFT model)


$$\text{objective}(\phi) = E_{(x,y) \sim D_{\pi_{\phi}^{\text{RL}}}} \left[r_{\theta}(x, y) - \beta \log \left(\pi_{\phi}^{\text{RL}}(y | x) / \pi^{\text{SFT}}(y | x) \right) \right] + \\ \gamma E_{x \sim D_{\text{pretrain}}} \left[\log(\pi_{\phi}^{\text{RL}}(x)) \right]$$

LLM training process


Step 3 – Alignment



- **Overoptimization problem of RM**
 - Larger RM is harder to hack
 - Gold: The true quality of the model output; Proxy: The score given by the RM used for model training.



(a) BoN



(b) RL

LLM training process

Step 3 – Alignment



- **Overoptimization problem of RM**

- Larger RM is harder to hack

- RM's scaling law


$$R_{\text{bon}}(d) = d(\alpha_{\text{bon}} - \beta_{\text{bon}}d)$$

$$R_{\text{RL}}(d) = d(\alpha_{\text{RL}} - \beta_{\text{RL}} \log d)$$


α : related to the size of RM parameters and data, indicating the positive gains from optimizing RM.

β : related to the size of RM parameters and data, indicating the negative impact of excessive optimization of RM.


d : The KL distance between the policy and the initial policy, which can be understood as the number of training steps.



(a) α_{bon}



(b) β_{bon}




(c) β_{RL}

LLM training process

Step 3 – Alignment



- **Overoptimization problem of RM**
 - Larger RM is harder to hack
 - At least 2K pieces of comparative data are required, the more the better



LLM training process

Step 3 – Alignment



- **Overoptimization problem of RM**
 - Multiple RM integrated/mixed RM data
 - Provide supervision signals from multiple perspectives to increase the difficulty of being hacked
 - For example, GPT-4 and Llama2 Chat both use multiple RMs

LLM training process

Step 3 – Alignment



- **Overoptimization problem of RM**
 - Fine-grained preference annotation
 - Process-supervised Reward Model (PRM)

The denominator of a fraction is 7 less than 3 times the numerator. If the fraction is equivalent to $\frac{2}{5}$, what is the numerator of the fraction? (Answer:)

Let's call the numerator x .

So the denominator is $3x-7$.

We know that $x/(3x-7) = 2/5$.

So $5x = 2(3x-7)$.

$5x = 6x - 14$.

So $x = 7$.



Section 5: General Language Model

- Model Architecture
- Challenges
- Mixed Precision Training



General Language Model (GLM)

Natural language: generation or comprehension?


- BERT, a self-supervised model: Utilizes bidirectional attention for text comprehension.
- GPT, a self-regressive model: Employs unidirectional attention for the generation of long texts.
- T5, an encoder-decoder model: Designed for encoding and decoding tasks, particularly useful for dialogue tasks.

Algorithm	NLU	Cond.Gen.	Uncond.Gen.
Self-regression	—	—	√
Self-supervised	√	×	×
Encoder-decoder	—	√	—
GLM	√	√	√



General Language Model (GLM)

- **Model Architecture: Autoregressive Blank Infilling**
- Pre-training: Bidirectional attention, predicting the masked portions
- Can freely set the proportions of three masking modes: word (MASK), sentence (sMASK), and document (gMASK).







General Language Model (GLM)

Model Architecture: UL2 - Hybrid Denoiser

- UL2: Highly similar pre-training structure objectives to GLM
- {X, R, S}-Denoiser is highly similar to GLM's MASK, sMASK, gMASK .
- By using UL2 training objectives for further training, U-PaLM significantly surpasses the performance of PaLM.







General Language Model (GLM)

Model Architecture: LayerNorm

- The original BERT paper adopted the Post-LN structure.
- The Post-LN structure is prone to divergence, and in recent years, models generally adopt the Pre-LN structure.
- In the training of several billion/multi-modal mixed-precision (FP16), Pre-LN is also unstable. A variant of Pre-LN, the Sandwich-LN structure, can alleviate this phenomenon.



Xiong, Ruibin, et al. "On layer normalization in the transformer architecture." International Conference on Machine Learning. PMLR, 2020

Ding, Ming, et al. "Cogview: Mastering text-to-image generation via transformers." Advances in Neural Information Processing Systems 34 (2021).

Devlin, Jacob, et al. "Bert: Pre-training of deep bidirectional transformers for language understanding." NAACL, 2019



General Language Model (GLM)

Model Architecture: LayerNorm

- DeepNet: Adjusting residuals, changing initialization → Stabilizing a thousand-layer Post-LN.

$$\text{DeepNorm}(x) = \text{LayerNorm}(\alpha \cdot x + g(x)), \quad \alpha > 1$$


- 130B scale experiment: DeepNorm is more stable than Sandwich-LN.
- Empirical observation: Post-LN is more sensitive to hyperparameters than Pre-LN and requires careful tuning.

```
def deepnorm(x):
    return LayerNorm(x *  $\alpha$  + f(x))

def deepnorm_init(w):
    if w is ['ffn', 'v_proj', 'out_proj']:
        nn.init.xavier_normal_(w, gain= $\beta$ )
    elif w is ['q_proj', 'k_proj']:
        nn.init.xavier_normal_(w, gain=1)
```

Architectures	Encoder		Decoder	
	α	β	α	β
Encoder-only (e.g., BERT)	$(2N)^{\frac{1}{4}}$	$(8N)^{-\frac{1}{4}}$	-	-
Decoder-only (e.g., GPT)	-	-	$(2M)^{\frac{1}{4}}$	$(8M)^{-\frac{1}{4}}$
Encoder-decoder (e.g., NMT, T5)	$0.81(N^4M)^{\frac{1}{16}}$	$0.87(N^4M)^{-\frac{1}{16}}$	$(3M)^{\frac{1}{4}}$	$(12M)^{-\frac{1}{4}}$

Gradient Norm



Deepnet: scaling transformers to 1,000 layers




General Language Model (GLM)

Model Architecture: Positional Embedding

- Positional encoding is a crucial module for Transformers to understand the order of the text, serving as a core component for processing long texts.
- Absolute Position Encoding: Three-angled, learnable.

$$\begin{cases} p_{k,2i} = \sin(k/10000^{2i/d}) \\ p_{k,2i+1} = \cos(k/10000^{2i/d}) \end{cases}$$



- Relative Position Encoding: Models the relative distances between words in attention.

$$q_i k_j^\top = (x_i + p_i) W_Q W_K^\top (x_j + p_j)^\top = (x_i W_Q + p_i W_Q) (W_K^\top x_j^\top + W_K^\top p_j^\top)$$

$$a_{i,j} = \text{softmax} \left(x_i W_Q \left(x_j W_K + \mathbf{R}_{i,j}^K \right)^\top \right)$$

Google^[3]

$$x_i W_Q W_K^\top x_j^\top + x_i W_Q W_K^\top \mathbf{R}_{i-j}^\top + \mathbf{u} W_Q W_K^\top x_j^\top + \mathbf{v} W_Q W_K^\top \mathbf{R}_{i-j}^\top$$

Transformer-XL^[4]



General Language Model (GLM)

Model Architecture: Positional Embedding


- ALiBi: Adding a bias matrix to Attention (different attention head coefficients)
 - Used by BLOOM and MPT.
 - Some extrapolation capability: MPT 64K extrapolated to 84K.
- RoPE: Rotation-based encoding, achieving relative encoding through absolute encoding.
 - Adopted by PaLM, GLM-130B, Llama, demonstrating stable performance improvement.
 - Simple implementation, efficient, and highly compatible (drawback: hardly extrapolates).

$q_1 \cdot k_1$				
$q_2 \cdot k_1$	$q_2 \cdot k_2$			
$q_3 \cdot k_1$	$q_3 \cdot k_2$	$q_3 \cdot k_3$		
$q_4 \cdot k_1$	$q_4 \cdot k_2$	$q_4 \cdot k_3$	$q_4 \cdot k_4$	
$q_5 \cdot k_1$	$q_5 \cdot k_2$	$q_5 \cdot k_3$	$q_5 \cdot k_4$	$q_5 \cdot k_5$

+

0				
-1	0			
-2	-1	0		
-3	-2	-1	0	
-4	-3	-2	-1	0

• m



$$(R_m q)^\top (R_n k) = q^\top R_m^\top R_n k = q^\top R_{n-m} k$$

$$R_{\Theta, m}^d = \begin{pmatrix} \cos m\theta_1 & -\sin m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ \sin m\theta_1 & \cos m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cos m\theta_2 & -\sin m\theta_2 & \cdots & 0 & 0 \\ 0 & 0 & \sin m\theta_2 & \cos m\theta_2 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \cos m\theta_{d/2} & -\sin m\theta_{d/2} \\ 0 & 0 & 0 & 0 & \cdots & \sin m\theta_{d/2} & \cos m\theta_{d/2} \end{pmatrix}$$



General Language Model (GLM)

Model Architecture: How to Achieve Longer Texts

- Classic Combination: RoPE + 2048 (Llama 1) / 4096 (Llama 2) Length Pre-training.
- How to achieve even longer texts (e.g., 32k)?
 - Directly continue pre-training on long texts => mediocre results.
 - Position interpolation + Continued pre-training on long texts => remarkably effective.
- NTK-aware Scaled RoPE: High-frequency extrapolation, low-frequency interpolation.
 - Effective extension from 2k without further pre-training to 8k.
 - Consideration from a beta-radix perspective.
 - Smooth transition from low-frequency fully equivalent interpolation to high-frequency fully equivalent extrapolation.


$$\left[\cos\left(\frac{n}{\beta^0}\right), \sin\left(\frac{n}{\beta^0}\right), \cos\left(\frac{n}{\beta^1}\right), \sin\left(\frac{n}{\beta^1}\right), \dots, \cos\left(\frac{n}{\beta^{d/2-1}}\right), \sin\left(\frac{n}{\beta^{d/2-1}}\right) \right]$$
$$\frac{n}{(\beta\lambda)^{d/2-1}} = \frac{n/k}{\beta^{d/2-1}} \quad \lambda = k^{2/(d-2)}$$



General Language Model (GLM)

Model Architecture: How to Achieve Longer Texts

- Perspective: Position interpolation + Continued training with Full Attention is already an extremely strong baseline.
- Long-text Evaluation: Perplexity (PPL) cannot truly reflect the model's performance on long texts.
 - MPT-30B-Chat-8K (1T 2k + 50B 8K) doesn't genuinely exhibit long-text capabilities.
 - LongBench: A multi-task long-text evaluation dataset in Chinese and English.





General Language Model (GLM)

Model Architecture: How to Achieve Longer Texts

- While extension methods are important, an efficient inference framework and downstream tasks are the core competitive factors.


Avg	Single-Doc QA	Multi-Doc QA	Summarization	Few-shot Learning	Code Completion	Synthetic Tasks	
GPT-3.5-Turbo-16k	44.0	39.8	38.7	26.5	67.1	54.1	37.8
Llama2-7B-chat-4k	31.0	24.9	22.6	24.7	60.0	48.1	5.9
LongChat-v1.5-7B-32k	34.3	28.7	20.6	26.7	60.0	54.1	15.8
XGen-7B-8k	28.3	24.6	20.4	24.7	56.2	38.6	5.3
InternLM-7B-8k	24.2	17.4	20.2	16.1	50.3	36.4	4.5
ChatGLM2-6B-32k	40.9	32.9	33.7	27.6	59.1	52.7	39.2
Vicuna-v1.5-7B-16k	31.9	28.0	18.6	26.0	66.2	47.3	5.5
ChatGLM3-6B-32k	48.5	40.3	46.6	29.5	68.1	56.2	50.5



General Language Model (GLM)

High Training Costs: The increase in model computation leads to a significant rise in overall training computational resources.

- GPT-3 used thousands of V100 GPUs, incurring a cost of \$4.6 million, with a total cost reportedly reaching \$12 million.
- OPT-176B training utilized 992 A100 GPUs, and PaLM 540B used 6144 TPUv4 units.






General Language Model (GLM)

Memory Challenges: The growth rate of the model's size far exceeds the growth rate of the computing device's memory.

- In addition to model parameters, the optimizer state occupies more GPU memory.





General Language Model (GLM)

Mixed Precision Training

- 16-bit floating-point numbers: FP16 / BF16
- Achieves 2-8 times performance improvement compared to FP32
- BF16: Supported on 3080, A100, TPU only

bfloat16: Brain Floating Point Format



Range: $\sim 1e^{-38}$ to $\sim 3e^{38}$

fp32: Single-precision IEEE Floating Point Format



Range: $\sim 1e^{-38}$ to $\sim 3e^{38}$

fp16: Half-precision IEEE Floating Point Format



Range: $\sim 5.96e^{-8}$ to 65504

NVIDIA A100 TENSOR CORE GPU SPECIFICATIONS
(SXM4 AND PCIE FORM FACTORS)

	A100 40GB PCIe	A100 80GB PCIe	A100 40GB SXM	A100 80GB SXM
FP64	9.7 TFLOPS			
FP64 Tensor Core	19.5 TFLOPS			
FP32	19.5 TFLOPS			
Tensor Float 32 (TF32)	156 TFLOPS 312 TFLOPS*			
BFLOAT16 Tensor Core	312 TFLOPS 624 TFLOPS*			
FP16 Tensor Core	312 TFLOPS 624 TFLOPS*			
INT8 Tensor Core	624 TOPS 1248 TOPS*			
GPU Memory	40GB HBM2	80GB HBM2e	40GB HBM2	80GB HBM2e
GPU Memory Bandwidth	1,555GB/s	1,935GB/s	1,555GB/s	2,039GB/s
Max Thermal Design Power (TDP)	250W	300W	400W	400W
Multi-Instance	Up to 7	Up to 7	Up to 7	Up to 7



General Language Model (GLM)

Mixed Precision Training: FP8

- The Hopper architecture introduces 8-bit floating-point numbers: FP8.
- **The Transformer Engine:** An acceleration library designed specifically for Transformers.
 - FP8 mixed precision support has been introduced on Hopper GPUs.
- Under the 16-bit mixed precision strategy, the emphasis is on converting GEMM to FP8 precision.

Model	Precision	Throughput (tok/sec)	TFLOPS	Speedup over NVIDIA A100 BF16
1B	bf16	43,352	394	2.2x
1B	fp8	53,721	489	2.7x
3B	bf16	22,987	412	2.2x
3B	fp8	29,313	525	2.8x
7B	fp8	13,433	580	3.0x
30B*	fp8	-	752	3.3x

规格	H800 SXM	H800 PCIe
FP64	1 TFLOP	0.8 TFLOP
FP64 Tensor Core	1 TFLOP	0.8 TFLOP
FP32	67 TFLOPS	51 TFLOPS
TF32 Tensor Core	989 TFLOPS*	756 TFLOPS*
BFLOAT16 Tensor Core	1979 TFLOPS*	1513 TFLOPS*
FP16 Tensor Core	1979 TFLOPS*	1513 TFLOPS*
FP8 Tensor Core	3958 TFLOPS*	3026 TFLOPS*
INT8 Tensor Core	3958 TOPS*	3026 TOPS*
GPU 显存	80GB	80GB
GPU 显存带宽	3.35TB/s	2TB/s
解码器	7 NVDEC 7 JPEG	7 NVDEC 7 JPEG
最大热设计功耗 (TDP)	最高 700 瓦 (可配置)	300 – 350 瓦 (可配置)
多实例 GPU	支持最大 7 个 MIG 实例, 每个实例 10 GB 显存	
外形规格	SXM	PCIe 双插槽风冷式
互连技术	NVLink™: 400GB/s PCIe 5.0: 128GB/s	NVLink: 400GB/s PCIe 5.0: 128GB/s
服务器选项	搭载 8 个 GPU 的 NVIDIA HGX™ H800 合作伙伴和 NVIDIA 认证系统和搭载 8 个 GPU 的 NVIDIA DGX™ H800 NVIDIA 认证系统 (NVIDIA-Certified Systems™)	搭载 1 至 8 个 GPU 的合作伙伴和 NVIDIA 认证系统
NVIDIA AI Enterprise 软件套件	附加	已包含




General Language Model (GLM)

Mixed Precision Training: FP8

- FP8-LM: Further reduces precision in communication, master weight, and optimizer state.
- The increase in TFLOPS mainly comes from the lower precision, saving memory and allowing for an increase in micro-batch size.

Low-bit Settings	Compute (GEMM)	Comm.	Master Weight	Optimizer States
FP32 #0	FP32	FP32	FP32	FP32+FP32
BF16 #1	<u>BF16</u>	FP32	FP32	FP32+FP32
FP8 #2	FP8	FP8	<u>FP16</u>	<u>FP8+FP16</u>
FP8 #3	FP8	FP8	<u>FP8</u>	FP8+FP16
FP8 #4	FP8	FP8	FP16	FP8+ <u>FP8</u>






General Language Model (GLM)

Data Parallelism:

- Data parallelism: Distributing data across multiple model instances.
- Distributed Stochastic Gradient Descent: The Parameter Server approach.
 - Load data from a distributed dataloader.
 - Perform forward + backward $\Rightarrow$ gradient.
 - Push gradient to PS (Parameter Server).
 - Pull weight from PS.
- Asynchronous: Typically, update once every N iterations.






General Language Model (GLM)

Data Parallelism:

- Disadvantages of Parameter Server:
 - Difficult to achieve load balancing.
- Distributed Stochastic Gradient Descent (The All-Reduce Approach)
 - Load data from distributed dataloader.
 - Perform forward + backward => gradient.
 - Perform All-Reduce.
 - Update weight locally.



```
import torch.nn.parallel as dist
from torch.nn.parallel import DistributedDataParallel as DDP

dist.init_process_group("nccl", rank=rank, world_size=world_size)
ddp_model = DDP(Model(), device_ids=[rank])

for batch in data_loader:
    loss = train_step(ddp_model, batch)
```



General Language Model (GLM)

Data Parallelism:

- Advantages:
 - Equal communication load for all nodes, achieving load balancing.
 - Easy to implement (PyTorch DDP).
- Disadvantages:
 - All workers perform one optimizer update.
 - Treating all nodes as homogeneous, not utilizing all bandwidth (CPU -> GPU).



General Language Model (GLM)

Data Parallelism: ZeRO Optimizer

- Disadvantage:
 - Full parameters, gradients, and optimizer states are stored on all nodes
- ZeRO: Zero Redundancy Optimizer
- Concept: Split optimizer states, gradients, and model weights
- FSDP (Fully Sharded Data Parallel) / HSDP

M is the number of parameters, N is the number of devices.


	Optimizer States (12M)	Gradients (2M)	Model Weights (2M)	Memory Cost	Communication Cost
Data Parallelism	Replicated	Replicated	Replicated	$16M$	all-reduce(2M)
ZeRO Stage 1	Partitioned	Replicated	Replicated	$4M + \frac{12M}{N}$	all-reduce(2M)
ZeRO Stage 2	Partitioned	Partitioned	Replicated	$2M + \frac{14M}{N}$	all-reduce(2M)
ZeRO Stage 3	Partitioned	Partitioned	Partitioned	$\frac{16M}{N}$	1.5 all-reduce(2M)




General Language Model (GLM)

Data Parallelism: Tensor Parallelism and Pipeline Parallelism

- Data parallelism cannot solve the problem when the model does not fit on a single GPU:
 - Full parameters, gradients, and optimizer states are stored on all nodes
- Tensor Parallelism: Split linear layers in the model across multiple GPUs and insert appropriate All-Reduce operators for data synchronization (high communication cost)
- Pipeline Parallelism: Divide the model into multiple stages and place them on different GPUs for computation



Tensor Parallelism



Pipeline Parallelism



The next lecture

- Lecture 15: Large Language Model II